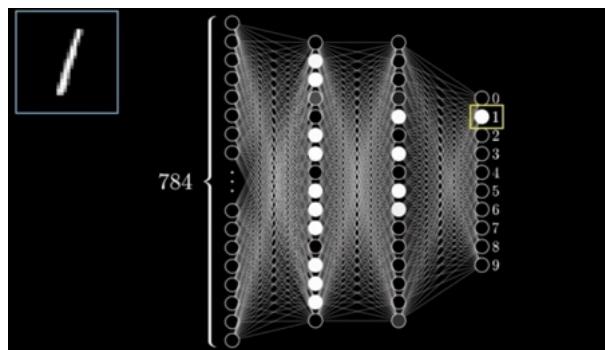
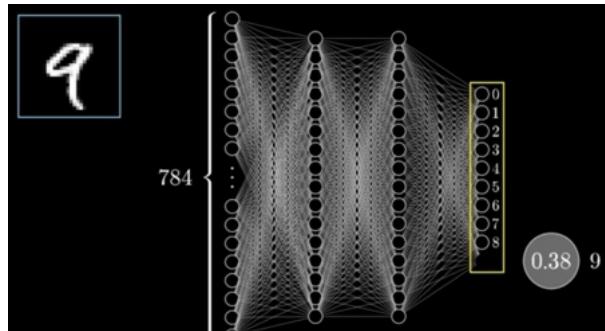


Neural Network:

Neuron:

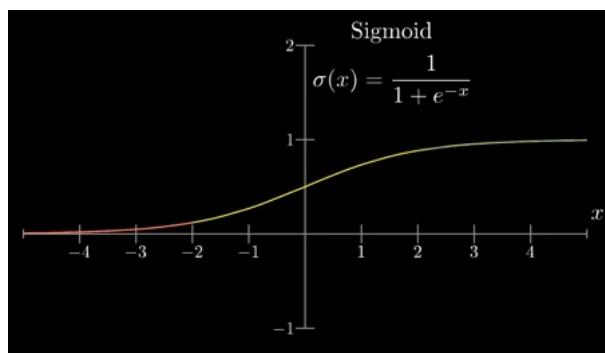
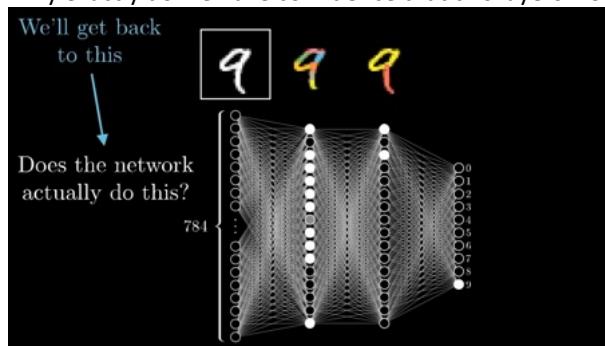
A thing that holds memory between 0 and 1

Take a 28*28 pixel image of a number now each cell holds a number between 0 to 1 (0 for dark 1 for white) this is known as activation

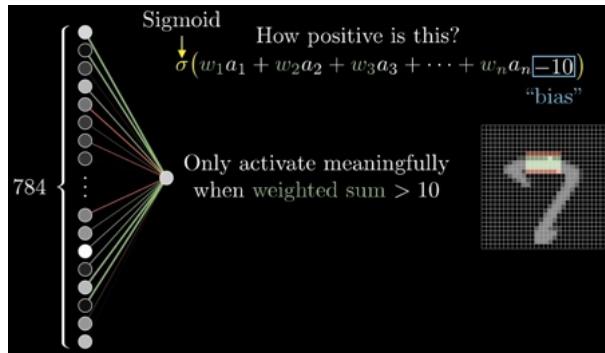


Activation of one layer determines the activation of the next layer and the answer is the brightest neuron in the final layer

Why exactly do we have confidence that this layers work??



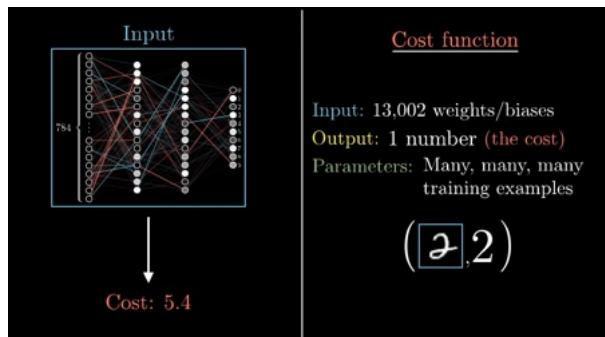
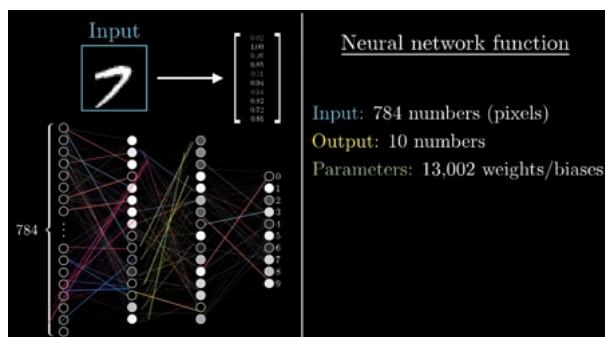
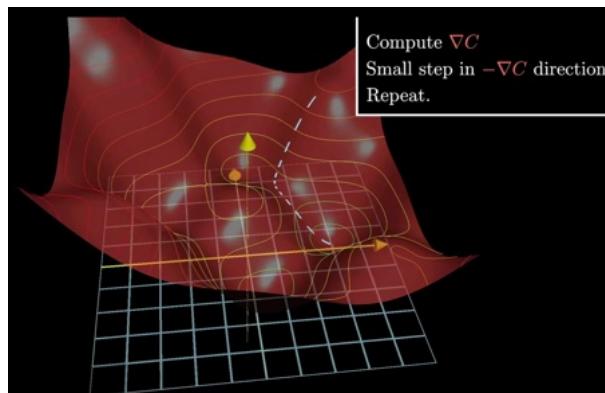
Now people are using ReLU because it is working increadably well for some deep neural networks



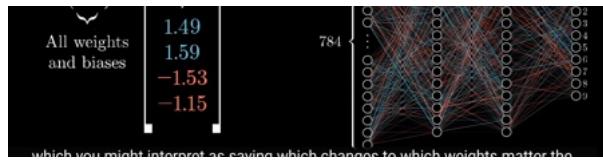
Cost function:

We should find the input(weights and bias) that minimizes this cost function

But it is very complicated for a traumtic function but we can choose something randomly and choose which direction to travel in (find the slope , if slope is >0 shift to the left if the slope<0 shift to the right)



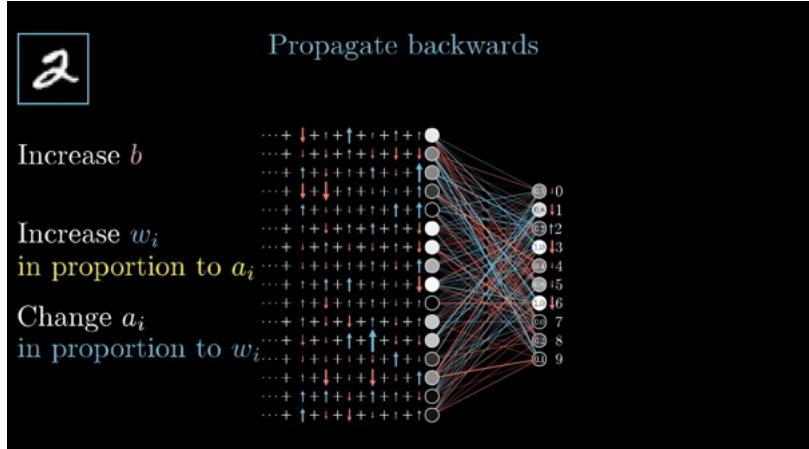
$-\nabla C(\dots) = \begin{bmatrix} 0.12 \\ 0.79 \\ -0.67 \\ 0.01 \\ \vdots \end{bmatrix}$
 Change by some small multiple of $-\nabla C(\dots)$



Despite identifying the number I can never draw them
And even if u give some nonsense it confidently says the answer for it

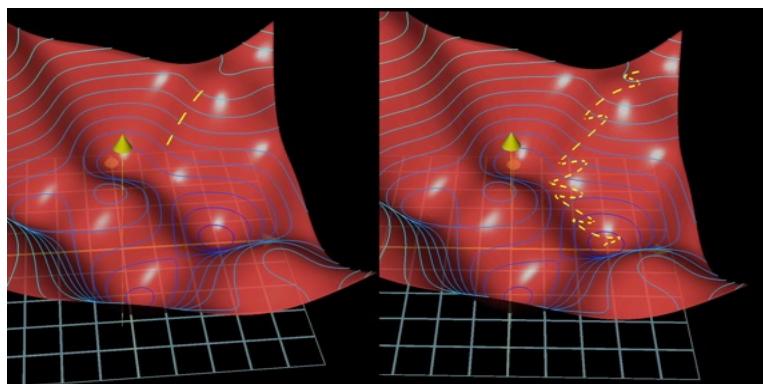
Back Propagation:

Is the algorithm how a single training data example will like to nudge the training weights and bias
A true gradient descent step will take the average for all the training data above and takes the decision



Stochastic gradient descent:

Dividing into small chunks and then rapidly stepping the down the hill with the help of the average of those few rowed data than doing for every data in the training set



MNIST data base

LLM:

A sophisticated mathematical function which predicts about what word comes next for any piece of text but instead of predicting one certain word it gives all possible words with their probabilities

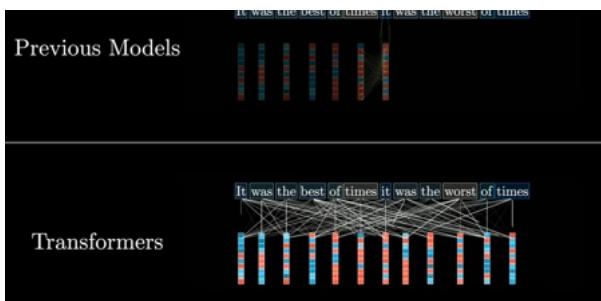
We have parameters in this which change continuously on training (same as how weights and bias are tuned based on the back propagation)

What all process does the chat bots undergo

- 1) pretraining
- 2) RLHF (Reinforcement learning with human feedback)

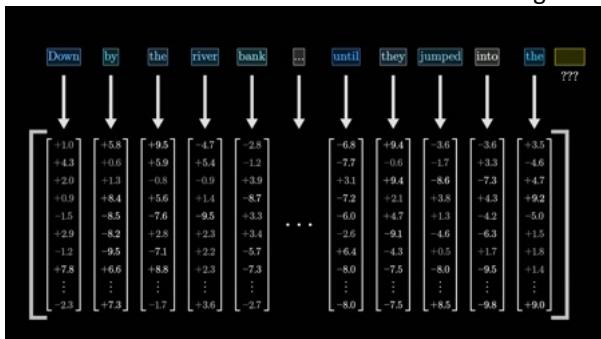
Before 2017 the words are analyzed one to one:



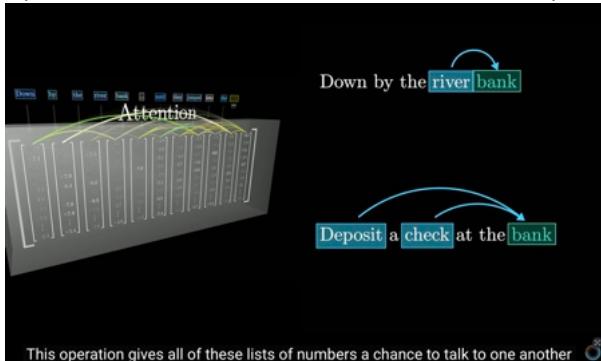


But transformer analyzes all the input text

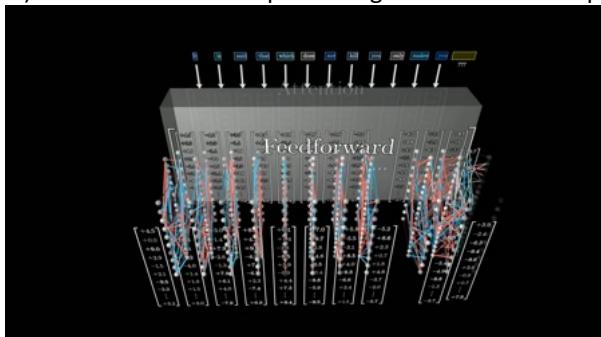
In transformer each word is associated with a large vector of numbers:



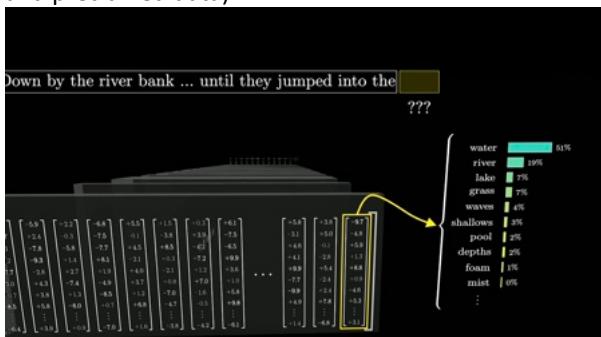
1) Attention: all the words talk with one another(which causes better context based meaning development for each word)



2) Feedforward: more processing of the words with pretrained data



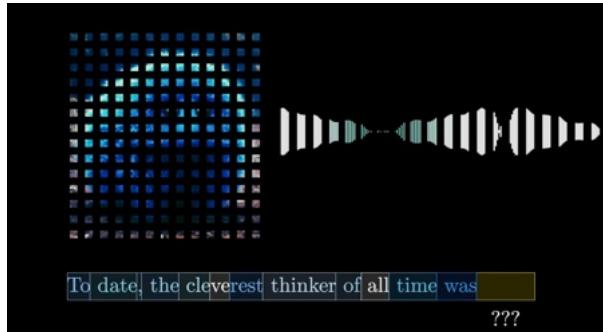
The last step produces a long vector with all the possible words along with their probabilities (this is influenced by our input and pretrained data)



GPT (Generative pretrained training)

Original transformer (in the publication of the Attention is all u need) is introduced for translating the text in one language to the other

Each little chunk is known as tokens:

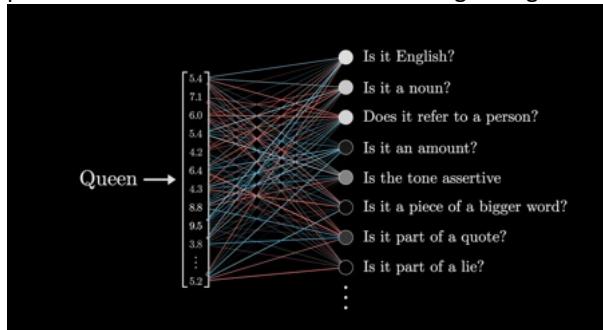


Each of these tokens is associated to a vector(as mentioned above)

If we think of these vectors in high dimensional coordinate space then words with similar meaning are near to eachother

Then these vectors are passed through the process attention

In multilayer perceptron/feed forward: Vectors don't talk to each other instead they all go through the same process in parallel to each other. It is more like asking a long list of questions to each vector and updating based on their response



Categories of weight

Total weights: 175,181,291,520	 GPT-3
Organized into 27,938 matrices	
Embedding	
Key	
Query	
Value	
Output	
Up-projection	
Down-projection	
Unembedding	